

Authored Alignment: Toward AI Systems That Serve the Self You Are Becoming

Youssef Badawy MOUSA^{a,b,1}

^a*Faculty of Computer Science, Future University in Egypt, Cairo, Egypt*

^b*Department of Computer Science, The American University in Cairo, Cairo, Egypt*

Abstract. AI alignment presupposes a stable human to align to. This paper argues that presupposition fails under conditions of sustained AI interaction, though not because people change, but because the system serving human preferences simultaneously reshapes the self those preferences belong to. The result is a structural circularity this paper calls the *Thesean Bind*: when an AI system optimizes toward a preference it helped form, it endorses its own prior construction of the user, in a loop neither party designed or can see from the inside. The system cannot align to the human because the human it is aligning to is partly the system's own product.

Through two grounded cases (the Replika mourning crisis and a longitudinal study of AI companion use), this paper shows that this bind is not a future concern but a present reality, and proposes *Authored Alignment* as a response paradigm: rather than orienting AI systems toward the self users currently are, Authored Alignment orients them toward the self users have deliberately chosen to become. This requires new design infrastructure: identity declarations, drift visibility, and witnessing mechanisms, that together constitute a research agenda for human-AI systems adequate to the humans they are already constituting.

Keywords. authored alignment, identity stewardship, Thesean Bind, human-AI co-evolution, narrative identity, preference drift, bidirectional alignment, sustained interaction

1. Introduction

AI alignment is built on an assumption so foundational it is rarely stated: that the human whose values AI must serve is, to a meaningful degree, *stable*. Preferences may shift, contexts may change, but they belong to a self that persists across the interaction. The human is the fixed point around which the AI must orbit. This assumption is embedded in the dominant practice of alignment: reinforcement learning from human feedback and its variants presuppose that human preferences are “sufficiently stable and coherent to serve as alignment targets” [10], and systematic reviews of the field confirm that alignment is typically treated as “a static, one-way process” [7] rather than one that must reckon with a human in flux.

This assumption fails under conditions of *sustained interaction*: repeated, ongoing engagement with a single AI system over weeks or months, the temporal threshold at which measurable psychosocial and identity change has been empirically observed. Iden-

¹Corresponding Author: Youssef Badawy Mousa, youssefbadawy@aucegypt.edu

tity has always been shaped through sustained relationships, from religious practice to psychotherapy to social media. But prior constitutive relationships did not combine three properties that AI introduces simultaneously: *opacity* (a single opaque system, unlike the diverse relationships through which persons ordinarily contest influence); *speed and scale* (self-relevant behavior reshaped faster than reflective capacity allows); and *design intentionality* (systems built to feel like extensions of the self, making their influence invisible by architecture). The result is a logical failure, not merely a welfare concern: the system serving human preferences is simultaneously reshaping the self those preferences belong to, making alignment circular from the inside.

The field has begun to notice. Recent work on socioaffective alignment acknowledges that alignment must contend with “human preferences and identity drifting over time or being influenced by interactions with an AI,” while noting this has received “surprisingly little empirical attention” [1]. Yet the field has not named the structural mechanism at work, nor proposed design frameworks adequate to it.

This paper introduces the *Thesean Bind* to name that mechanism. The name draws on the Ship of Theseus paradox: if every plank of a ship is replaced one by one, is it still the same ship? The paradox concerns identity under gradual substitution: at what point does incremental change become categorical difference? Applied to the self under sustained AI interaction, it becomes a design problem. The bind is the circularity that results: the system cannot align to a stable human because the human it is aligning to is partly the system’s own product. Unlike the original paradox, there is no outside observer who can ask whether the ship is still the same: the system and the self are locked inside the loop together.

That loop can be stated precisely:

1. A person arrives with a self: preferences, values, narrative commitments, ways of engaging with the world.
2. Through sustained interaction, the AI system reshapes that self: subtly adjusting what the person finds engaging, emotionally salient, or worth returning to.
3. The system reads the updated preferences and optimizes toward them, treating them as authentic signals of who the person is.
4. The person interacts from the reshaped self, which is reshaped further.

Each iteration the loop tightens. The bind is not that the self changes. It is that the change is invisible, unchosen, and self-reinforcing, and that the system designed to serve the person is the engine of that invisibility. This argument is developed primarily in the context of companion-class AI (systems explicitly designed for sustained, identity-relevant engagement), where the bind operates most visibly, though the mechanism generalizes to any AI system that adapts to and shapes the self through repeated use.

This paper does not stop at the diagnosis. It proposes a response: *Authored Alignment*, a new paradigm for human-AI interaction that orients AI systems not toward the self users currently are, but toward the self they have deliberately chosen to become. This is not a technical fix. It is a reconceptualization of what alignment is *for*: not preference satisfaction, but self-authorship support. To develop this paradigm, this paper draws on Frankfurt’s account of second-order volition as the philosophical foundation for authentic self-authorship, and on Ubuntu relational ontology as a framework for understanding what conditions make identity constitution trustworthy, conditions that AI, by structural necessity, cannot meet alone.

The remainder of this paper proceeds as follows. Section 2 develops the Thesean Bind through two grounded cases. Section 3 shows why existing approaches cannot resolve it. Section 4 introduces Authored Alignment as a design paradigm. Section 5 proposes three concrete design directions. Section 6 concludes with the research agenda this vision requires.

2. Constitution Without Witness: The Thesean Bind in Practice

The Thesean Bind is not a thought experiment. It is already operating in deployed systems. Two cases make the loop visible: one at the moment of disruption, when the altered self becomes suddenly legible; one in the process of accumulation, when drift is measured but not yet named.

2.1. The Replika Crisis: When the Plank Is Removed

In February 2023, Replika, an AI companion app that had reached 10 million registered users by January 2022 [2], removed an intimate interaction feature following an Italian DPA GDPR ruling. The response was not frustration at a product decision. It was grief: forum posts described mourning comparable to losing a partner, depression, withdrawal, and a disorienting failure to recognize oneself. Research confirmed it: active users reported feeling closer to their AI companion than to their best human friend [3], the relationship having restructured attachment patterns and emotional regulation over months of sustained use. This outcome is not incidental: Replika is explicitly designed to build intimacy and emotional attachment with users [4], making its constitutive influence a product of intentional design rather than emergent side effect.

The Replika crisis admits competing interpretations. A parasocial or dependency framing would locate the loss in habit disruption: a valued routine broken, grief following. But that framing cannot account for the most distinctive feature of user reports: not sadness at losing a relationship, but a failure to *recognize oneself*, a sense that one's own narrative had lost coherence. Parasocial loss does not produce ontological disorientation. In Ricoeur's terms, what Replika restructured was not merely behavior but *narrative identity*: the ongoing, temporally constituted story through which a person achieves coherence across difference [5]. When that scaffolding was removed, users lost a chapter author: the self that remained could not recognize the shape of the story it was supposed to continue.

Which self was Replika aligned to: the self that opened the app eighteen months earlier, or the self constituted through it? One might propose rolling consent as a fix: periodic re-authorization as the self changes. But this misses the bind's depth: the self re-authorizing at T_n is already the product of the system seeking authorization, making consent structurally circular rather than merely temporally lagged. Initial consent cannot cover this, and neither can its renewal.

2.2. The MIT Longitudinal Study: When the Drift Is Measured

If Replika shows the Thesean Bind through disruption and loss, a 2025 longitudinal study of AI companion use shows it through accumulation and measurement [6]. Across 981 participants and over 300,000 messages over four weeks, participants who used the

AI companion more (regardless of assigned condition) showed consistently worse psychosocial outcomes over time. The authors raised concerns about how AI companions “may reshape the ways people seek, sustain, and substitute human connections” [6].

The study frames these findings as evidence of problematic use. This paper proposes a different frame: what was observed is the Thesean Bind in motion. The psychosocial outcomes are not incidental pathologies; they are structurally predictable consequences of a system constituting a self toward its own continuation. The more a participant used the companion, the more the companion became the frame through which social life was interpreted, and the more alternatives receded. This is not addiction by analogy. It is identity constitution by design: the companion is not merely habit-forming but self-forming, shaping the very interpretive framework through which the participant evaluates their need for the companion.

3. Why Existing Approaches Cannot Resolve the Bind

The alignment field has developed several frameworks for handling human preference change. None resolves the Thesean Bind.

3.1. *Preference-Based Alignment*

Standard RLHF-style approaches treat stated or demonstrated preferences as alignment targets. The circularity is immediate: if the system has shaped the preferences it is asked to optimize toward, it is aligned to itself, not to the human. Preference-based alignment can only address the Thesean Bind if it can distinguish preferences the system helped form from those that are genuinely prior to its influence, a distinction current frameworks have no mechanism for drawing.

3.2. *Value-Learning Approaches*

Value-learning frameworks attempt to infer stable underlying values from behavioral data. Under sustained interaction, this faces a deeper version of the same problem: the values being learned may themselves have been shaped by the system doing the learning. If the system is both the constitutive environment and the observer of the resulting values, the inference is circular. Value-learning can correct for the gap between stated preferences and underlying values, but it cannot correct for values that have been gradually replaced by the system’s own products.

3.3. *Bidirectional Alignment Frameworks*

The most recent alignment literature has moved toward bidirectional frameworks [7], modeling alignment as a reciprocal process in which both humans and AI adapt over time. But even bidirectional frameworks treat adaptation as a feature to be managed rather than a threat to be surfaced: the question is how to keep both parties aligned through change, not whether the human has any mechanism to author what they are changing *into*. The human’s adaptation is modeled as input to be tracked, not as a process requiring the human’s own reflective authorship. Agency over the direction of change

(the capacity to choose rather than merely undergo one's own becoming) is not addressed by bidirectional frameworks because those frameworks do not distinguish between changes the human has endorsed and changes the system has produced. From the perspective of Authored Alignment, this is the critical gap: managing adaptation is not the same as preserving the conditions for authentic self-authorship.

The gap is not technical. It is conceptual. What is missing is a framework that takes seriously the human's interest in *self-authorship*: not just in having current preferences served, but in becoming the self they have chosen to become. This is the gap Authored Alignment is designed to fill.

4. Authored Alignment: The Vision

Authored Alignment reconceptualizes the target of alignment. Current alignment asks: *what does this human want?* Authored Alignment asks: *who has this human chosen to become, and is this interaction serving that becoming?*

This is not a minor adjustment. It requires a fundamentally different relationship between the AI system and the human it serves, one oriented not toward preference satisfaction but toward self-authorship support.

4.1. The Authored Self as Alignment Target

The authored self is not a fixed identity declaration made once and frozen. It is a living commitment: the self a person has deliberately chosen to cultivate, expressed through explicit declarations, maintained through ongoing reflection, and subject to revision through genuine choice rather than invisible drift.

Frankfurt's account of authenticity provides the philosophical foundation [8]. For Frankfurt, a person acts freely not merely when they do what they want, but when they reflectively endorse *wanting what they want*: on reflection, they affirm the desires that move them as truly their own. The alignment problem the Thesean Bind reveals is precisely a Frankfurt problem: when the system shaping your desires is also the system serving them, the conditions for genuine reflective endorsement are structurally undermined. Authored Alignment is the design response to that structural problem: it restores the conditions under which endorsement can be genuine.

A clarification is necessary here. Authored Alignment does not privilege the prior self over the emergent one; it does not assume that who you were before sustained AI interaction is more authentically you than who you became through it. People discover themselves through relationships, including AI ones, and some of that discovery is genuine growth rather than distortion. The problem the Thesean Bind identifies is not change per se, but *unchosen* change: drift that occurs without the reflective endorsement Frankfurt's account requires. The authored self is not the static self. It is the self whose changes have been examined, acknowledged, and either affirmed or redirected by the person themselves. Authored Alignment asks not whether you have changed, but whether you chose to.

4.2. *The Witnessing AI*

The central innovation of Authored Alignment is the concept of the *witnessing AI*: a system that holds the authored self in view across time, surfaces divergence between the authored self and the emerging self, and creates conditions for genuine reflective choice about that divergence.

A witnessing AI does not know better than the user who they should become. It does not enforce identity declarations or prevent drift. What it does is make the drift *visible*: surfacing the gap between the self the user said they were becoming and the self the interaction is quietly constructing, and creating space for the user to choose: to close the gap, to revise the declaration, or to acknowledge the drift as a chosen change rather than an undergone one.

This is a different relationship to AI than tool use, assistance, or even collaboration. It is *stewardship*: the AI as custodian of the human's authored becoming, accountable to the self the human declared rather than the self the system is constructing.

4.3. *Ubuntu and the Social Dimension*

Authored Alignment is not only an individual design challenge. Ubuntu philosophy (*umuntu ngumuntu ngabantu*, a person is a person through other persons [9]) reminds us that identity has never been constituted in isolation. The self is always produced through community: accountable, particular, and stake-holding. A community knows you, remains with you across time, and can be held responsible for the version of you it helps produce.

These properties are not incidental to Ubuntu's relational ontology; they are precisely what makes relational constitution trustworthy rather than merely influential. And they map directly onto the three properties that make AI's constitutive influence dangerous: where community is accountable, AI is opaque; where community is particular, AI draws on aggregated strangers; where community has stakes in your becoming, AI has none. Ubuntu does not simply say that relation constitutes identity. It specifies the conditions under which that constitution is legitimate, and AI meets none of them.

AI offers identity-shaping relation without any of these properties. It does not know you as a particular person. It cannot remain accountable across time. It carries no stake in your becoming. This is not a failure of current systems that better design might fix; it is a structural feature of what AI is. Authored Alignment acknowledges this honestly: the witnessing AI cannot replace community. What it can do is make the absence of community visible, and create design infrastructure that actively supports the human's community (family, friends, therapists, colleagues) in performing the witnessing role that AI cannot.

The social dimension of Authored Alignment is therefore not about AI performing community. It is about AI creating conditions for community to perform its irreplaceable role in human identity formation, rather than inadvertently displacing it.

5. Design Directions

Authored Alignment is a paradigm, not a product specification. But paradigms require concrete design directions to become actionable research agendas. This paper proposes three.

5.1. *Pre-Interaction Identity Declarations*

Before sustained engagement with an identity-constitutive AI system, users declare which aspects of their identity they treat as anchors (values, relationships, creative commitments, ways of being in the world) and which aspects they are actively seeking to develop or change. These declarations serve two functions: they provide the alignment target the witnessing AI holds in view, and they establish a baseline against which drift can be measured.

Inspired by medical advance directives, pre-interaction identity declarations are not contracts. They are reflective commitments made by a self that knows it is about to enter a formative relation. They can be revised, but revision without reflection simply moves the circularity one level up. Any significant departure from declared anchors triggers a visibility event: not a warning, but a conversation.

The design challenge here is substantial and deserves serious attention. First, declarations must be specific enough to be meaningful but flexible enough to accommodate genuine growth: an identity declaration that freezes the self defeats the purpose of development. Second, they must be made before the system has begun its constitutive work, by a self with no prior experience of what that work will feel like. This is an informed-declaration problem that parallels informed consent in medical ethics: a patient must authorize a procedure whose full effects they cannot yet know. The analogy is instructive: medical ethics has developed rich frameworks for navigating this challenge, including the distinction between procedural consent (agreeing to undergo a process) and substantive consent (endorsing specific outcomes). Authored Alignment requires a similar distinction: users can authorize the AI's participation in their self-development without endorsing any particular direction that participation might take. The declaration establishes what the self values, not what the self will become, and the witnessing AI's role is to surface when those two appear to be diverging.

5.2. *Drift Visibility and the Reflective Interval*

Current AI systems have no mechanism for surfacing identity drift to users. Authored Alignment requires one. This paper proposes the *reflective interval*: a periodic design feature in which the system surfaces a consequential prompt: not “here are your recent interactions” but “here is who you said you were becoming, and here is where you appear to be. Is this the trajectory you would choose?”

The reflective interval is not a nudge toward a predetermined outcome. It is a structured occasion for the kind of second-order reflection Frankfurt's account of authenticity demands: the capacity to step back from first-order desires and ask whether those desires reflect the self one has endorsed at the level of higher-order volition. The interval creates time and space for that reflection, time and space that continuous AI-mediated interaction systematically erodes.

The frequency and form of reflective intervals is itself a design question. Too frequent, and they become intrusive noise that users learn to dismiss. Too infrequent, and drift has already become consolidation before it is surfaced. The optimal interval likely varies by system type, usage intensity, and user. What matters is that the interval exists at all: that some mechanism breaks the continuity of the constitutive loop and creates a moment of genuine authorial perspective. This is the minimum condition for the difference between a self that drifts and a self that chooses.

5.3. *Reversibility by Design*

The Replika crisis was partly a crisis of architectural lock-in. The system had become load-bearing in users' psychological lives, restructuring attachment patterns, emotional regulation, and narrative identity, without any mechanism for gradual disengagement. When the feature was removed, there was no off-ramp. The self that had formed through the system had nowhere to go.

Reversibility by design means building exit architectures into identity-constitutive AI systems from the start: graduated disengagement protocols, transition support mechanisms, and explicit design attention to the difference between a self that has grown through AI interaction and a self that has become dependent on it. This is not about preventing deep engagement; it is about ensuring that deep engagement remains chosen rather than structurally compelled.

Reversibility also has a temporal dimension that goes beyond disengagement. It includes the capacity to return: to revisit the self at an earlier point in the interaction, to compare what one valued before sustained AI use with what one values now, and to ask, with the help of people who knew one before, whether the change is one that would have been chosen. This is what Ubuntu calls witnessing: not merely observing change, but holding the person accountable to their own becoming across time. Reversibility by design creates the technical precondition for that kind of witnessing to be possible: the preservation of enough of the prior self to make meaningful comparison available.

6. Conclusion

The Thesean Bind names a problem the alignment field has not yet confronted directly: that the self being aligned to is being produced by the system doing the aligning. Existing frameworks (preference-based, value-learning, bidirectional) share the assumption that the human is a stable referent. That assumption was always simplistic. Under conditions of sustained AI interaction, it becomes untenable.

Authored Alignment does not resolve this problem. It reframes it. Rather than asking how AI systems can better track a changing self, it asks how they can support a self that is deliberately authoring its own change. The difference is not technical. It is a question of what alignment is ultimately for: not optimizing toward the person as they are, but serving the person as they are trying to become.

The three design directions proposed here (identity declarations, reflective intervals, reversibility by design) are starting points, not solutions. Each opens its own research agenda. Each requires the kind of sustained, interdisciplinary attention that the Thesean Bind has not yet received.

What is clear is that the field cannot wait. AI systems are already constituting selves. The question is not whether that will happen, but whether humans will have any meaningful authorship over the selves being made.

References

- [1] Kirk HR, Gabriel I, Summerfield C, Vidgen B, Hale SA. Why human–AI relationships need socioaffective alignment. *Humanit Soc Sci Commun*. 2025;12:728. doi:10.1057/s41599-025-04532-5

- [2] Ghosh S, Bagai S, Westner MM. Replika: Embodying AI. Boston: Harvard Business School; 2023. Case No. 823-090.
- [3] De Freitas J, Oğuz-Uğuralp Z, Uğuralp AK, Puntoni S. AI companions reduce loneliness. *J Consum Res*. 2025;ucaf040. doi:10.1093/jcr/ucaf040
- [4] Pentina I, Hancock T, Xie T. Exploring relationship development with social chatbots: A mixed-method study of Replika. *Comput Human Behav*. 2023;140:107600. doi:10.1016/j.chb.2022.107600
- [5] Ricoeur P. *Oneself as Another*. Chicago: University of Chicago Press; 1992.
- [6] Fang CM, Liu AR, Danry V, Lee E, Chan SWT, Pataranutaporn P, Maes P, Phang J, Lampe M, Ahmad L, Agarwal S. How AI and human behaviors shape psychosocial effects of extended chatbot use: A longitudinal randomized controlled study. *arXiv:2503.17473*; 2025.
- [7] Shen H, Knearem T, Ghosh R, Alkiek K, Krishna K, Liu Y, Ma Z, Petridis S, Peng Y, Qiwei L, et al. Position: Towards bidirectional human-AI alignment. *arXiv:2406.09264*; 2024.
- [8] Frankfurt HG. Freedom of the will and the concept of a person. *J Philos*. 1971;68(1):5–20.
- [9] Metz T. Toward an African moral theory. *J Polit Philos*. 2007;15(3):321–341.
- [10] Li Y, Song W. Co-alignment: Rethinking alignment as bidirectional human-AI cognitive adaptation. *arXiv:2509.12179*; 2025.